

# Efficient Pre-trained Features and Recurrent Pseudo-Labeling in Unsupervised Domain Adaptation

Youshan Zhang      Brian D. Davison

Computer Science and Engineering, Lehigh University, Bethlehem, PA, USA  
{yoz217, bdd3}@lehigh.edu

## Abstract

*Domain adaptation (DA) mitigates the domain shift problem when transferring knowledge from one annotated domain to another similar but different unlabeled domain. However, existing models often utilize one of the ImageNet models as the backbone without exploring others, and fine-tuning or retraining the backbone ImageNet model is also time-consuming. Moreover, pseudo-labeling has been used to improve the performance in the target domain, while how to generate confident pseudo labels and explicitly align domain distributions has not been well addressed. In this paper, we show how to efficiently opt for the best pre-trained features from seventeen well-known ImageNet models in unsupervised DA problems. In addition, we propose a recurrent pseudo-labeling model using the best pre-trained features (termed PRPL) to improve classification performance. To show the effectiveness of PRPL, we evaluate it on three benchmark datasets, Office+Caltech-10, Office-31, and Office-Home. Extensive experiments show that our model reduces computation time and boosts the mean accuracy to 98.1%, 92.4%, and 81.2%, respectively, substantially outperforming the state of the art.*

## 1. Introduction

With the explosive growth of information in the current era, there are massive amounts of data from multiple sources and corresponding to varied scenarios. However, not all tasks have enough annotated data for training, and collecting sufficient labeled data is a big investment of time and effort. Therefore, in order to build machine learning models it is often necessary to transfer knowledge from one labeled domain to an unlabeled domain. Due to dataset bias or domain shift [17], the generalization ability of the learned model on the unlabeled domain has been severely compromised. Domain adaptation (DA) is proposed to circumvent the domain shift problem.

Unsupervised domain adaptation (UDA) transfers knowl-

edge learned from a label-rich source domain to a fully unlabeled target domain [16]. Most prior methods focus on matching (marginal, conditional, and joint) distributions between two domains to learn domain-invariant representations. Maximum Mean Discrepancy (MMD) is one of the most popular distance metrics when minimizing differences between two distributions. Long et al. [14] proposed a Deep Adaptation Network (DAN) that considered multiple kernels of MMD functions. Recently, Kang et al. [11] extended MMD to the contrastive domain discrepancy loss. However, these distance-based metrics can also mix samples of different classes together. Recently, adversarial learning has shown its power in learning domain invariant representations. The domain discriminator aims to distinguish the source domain from the target domain, while the feature extractor aims to learn domain-invariant representations to fool the domain discriminator [11]. Sometimes, pseudo-labeling is proposed to learn the target discriminative representations [31, 32]. However, the credibility of these pseudo labels is unknown.

To address the above challenges, this paper provides two specific contributions:

1. To reduce computation time, we extract features from seventeen pre-trained ImageNet models and then design a fast and efficient unsupervised metric to select the best pre-trained features for the domain transfer tasks.
2. We develop a recurrent pseudo-labeling paradigm to continuously select high confidence transfer examples from the target domain and minimize the marginal and conditional discrepancies between the two domains.

We conduct extensive experiments on three benchmark datasets (Office + Caltech-10, Office-31, and Office-Home), achieving higher accuracy than state-of-the-art methods.

## 2. Related work

**Pre-training** Pre-training is one of the dominant components of transfer learning. Recent deep networks often apply a pre-trained network (typically trained on the ImageNet dataset) as the initialization for object recognition and segmentation. As in many computer vision tasks, it is often



















